

Review of Computer Architetcure

A Sahu
Deptt. of Comp. Sc. & Engg.
IIT Guwahati

Outline

- Computer organization Vs Architecture
- Processor architecture
- Pipeline architecture
 - Data, resource and branch hazards
- Superscalar & VLIW architecture
- Memory hierarchy
- Reference

Computer Organization vs Architecture

Comp Organization => Digital Logic Module
Logic and Low level

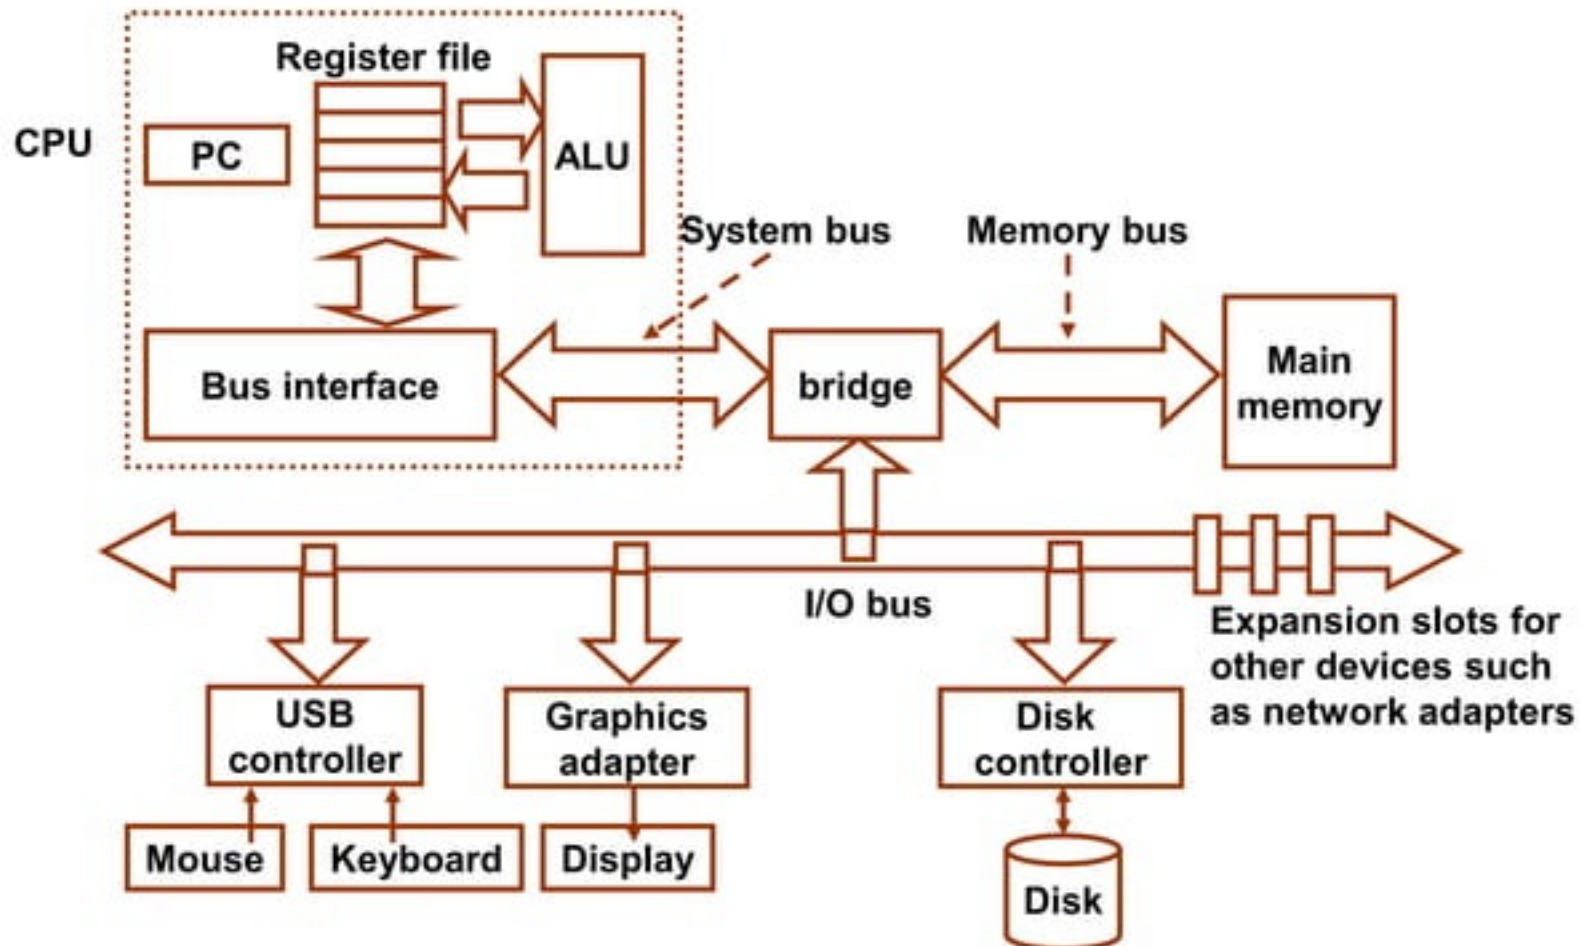
=====

Comp Architecture = > ISA Design, MicroArch
Design

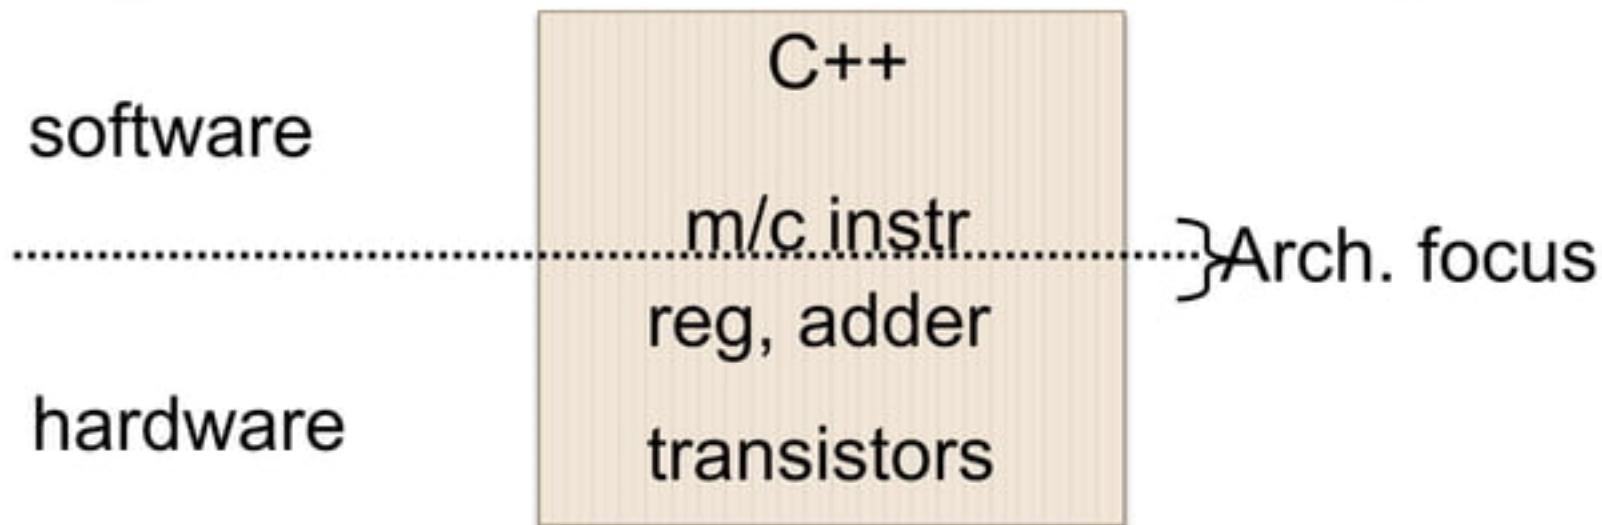
Algorithm for

- Designing best micro architecture,
- Pipeline model,
- Branch prediction strategy, memory management

Hardware abstraction



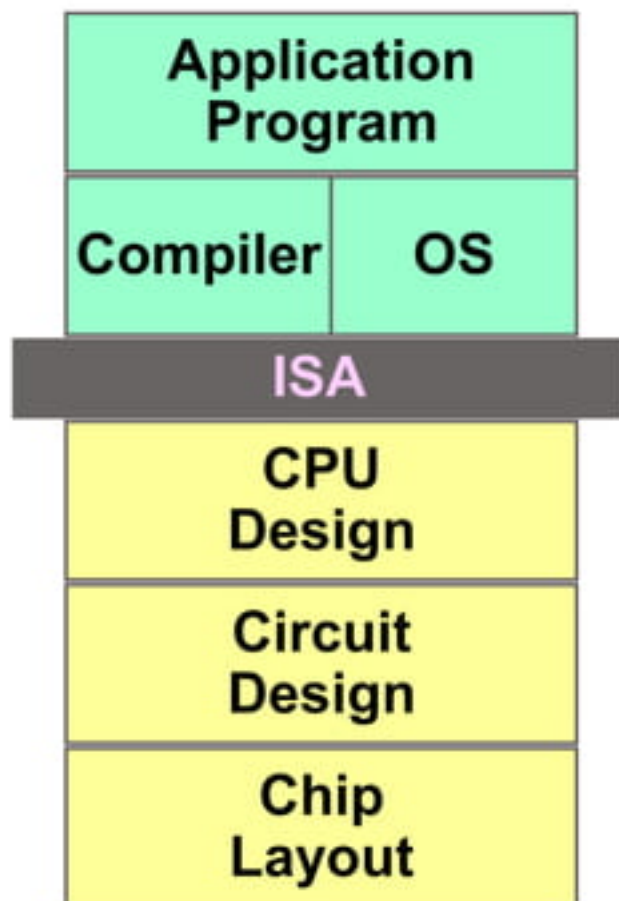
Hardware/software interface



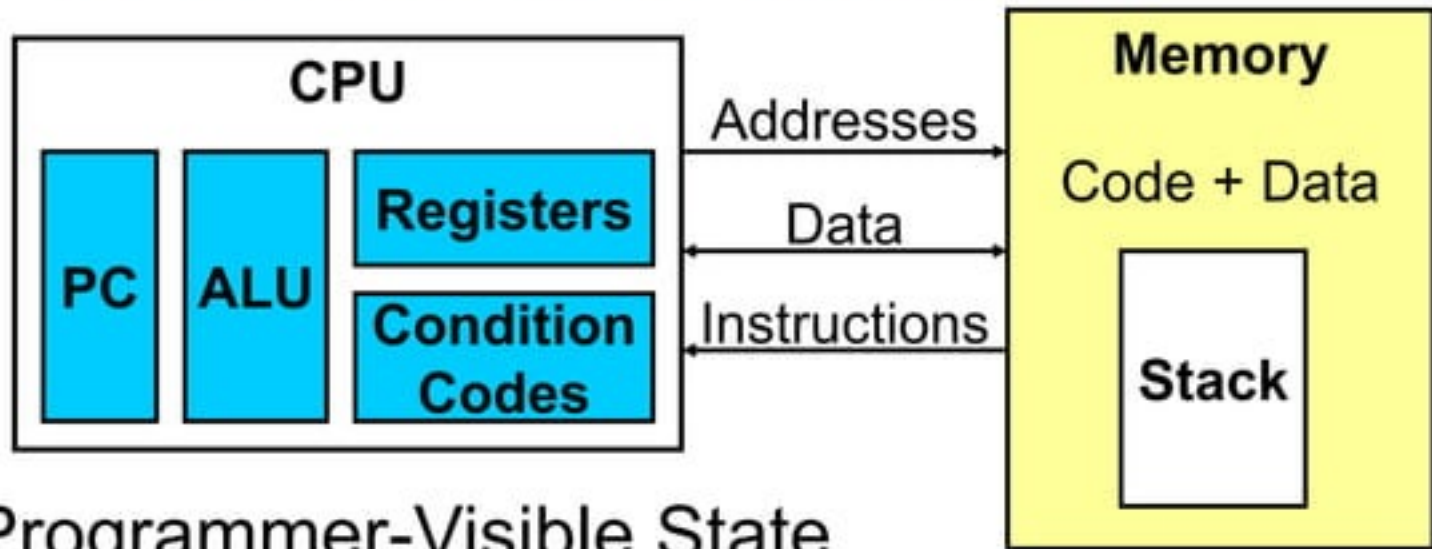
- Instruction set architecture
 - Lowest level visible to a programmer
- Micro architecture
 - Fills the gap between instructions and logic modules

Instruction Set Architecture

- Assembly Language View
 - Processor state (RF, mem)
 - Instruction set and encoding
- Layer of Abstraction
 - Above: how to program machine - HLL, OS
 - Below: what needs to be



The Abstract Machine



- **Programmer-Visible State**

- PC Program Counter
- Register File
 - heavily used data
- Condition Codes

- **Memory**

- Byte array
- Code + data
- stack

Instructions

- **Language of Machine**
- **Easily interpreted**
- **primitive compared to HLLs**

- **Instruction set design goals**
 - **maximize performance,**
 - **minimize cost,**
 - **reduce design time**

Instructions

- All MIPS Instructions: 32 bit long, have 3 operands
- Operand order is fixed (destination first)

Example:

C code: $A = B + C$

MIPS code: add \$s0, \$s1, \$s2

(associated with variables by compiler)

- Registers numbers 0 .. 31, e.g.,
\$t0=8,\$t1=9,\$s0=16,\$s1=17 etc.

- 000000 10001 10010 01000 00000

Instructions LD/ST & Control

- Load and store instructions
- Example:

C code: $A[8] = h + A[8];$

MIPS code: lw \$t0, 32(\$s3)
add \$t0, \$s2, \$t0
sw \$t0, 32(\$s3)

- Example: lw \$t0, 32(\$s2)

35	18	9	32
op	rs	rt	16 bit number

- Example:

if (i != j)	beq \$s4, \$s5, Lab1
h = i + j;	add \$s3, \$s4, \$s5
else	j Lab2
h = i - j;	Lab1: sub \$s3, \$s4, \$s5
	Lab2: ...

What constitutes ISA?

- **Set of basic/primitive operations**
 - Arithmetic, Logical, Relational, Branch/jump, Data movement
- **Storage structure – registers/memory**
 - Register-less machine, ACC based machine, A few special purpose registers, Several Gen purpose registers, Large number of registers
- **How addresses are specified**
 - Direct, Indirect, Base vs. Index, Auto incr and auto decr, Pre (post) incr/decr, Stack
- **How operand are specified**
 - 3 address machine $r1 = r2 + r3$, 2 address machine $r1 = r1 + r2$
 - 1 address machine $Acc = Acc + x$ (Acc is implicit)

RISC vs. CISC

- **RISC**
 - Uniformity of instructions,
 - Simple set of operations and addressing modes,
 - Register based architecture with 3 address instructions
- **RISC: Virtually all new ISA since 1982**
 - ARM, MIPS, SPARC, HP's PA-RISC, PowerPC, Alpha, CDC 6600
- **CISC** : Minimize code size, make assembly language easy

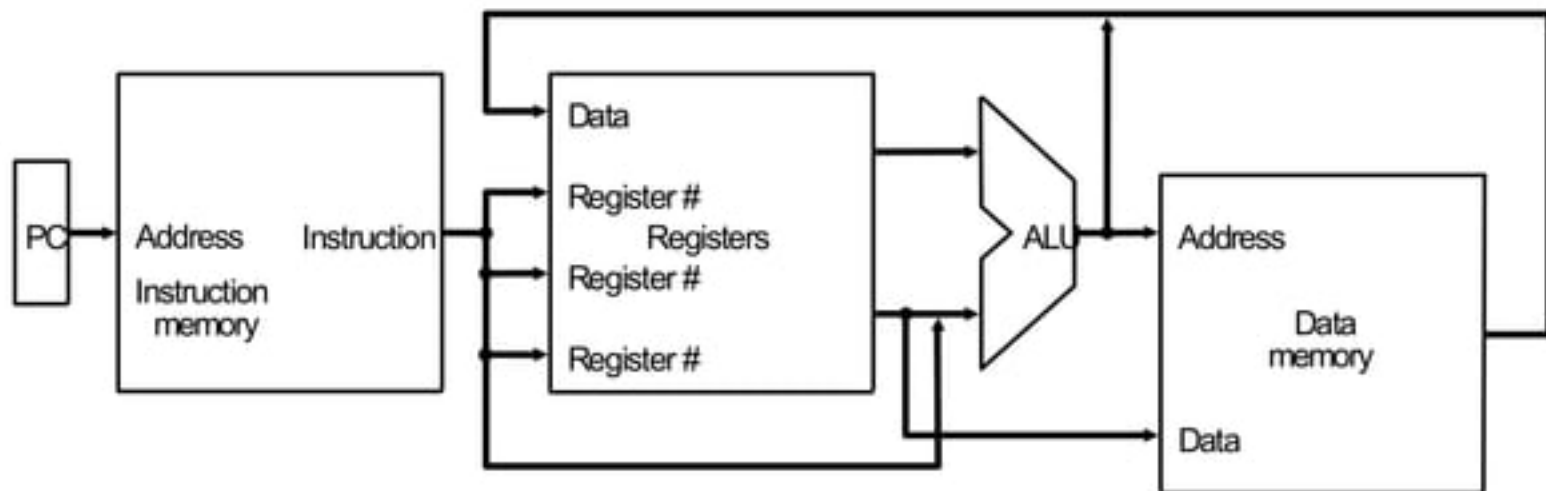
MIPS subset for implementation

- Arithmetic - logic instructions
 - add, sub, and, or, slt
- Memory reference instructions
 - lw, sw
- Control flow instructions
 - beq, j


Incremental changes in the design to include other instructions will be discussed later

Design overview

- Use the program counter (PC) to supply instruction address
- Get the instruction from memory
- Read registers
- Use the instruction to decide exactly what to do



Division into data path and control



Building block types


Two types of functional units:

- elements that operate on data values (combinational)
 - output is function of current input, no memory
 - Examples
 - gates: and, or, nand, nor, xor, inverter, Multiplexer, decoder, adder, subtractor, comparator, ALU, array multipliers
- elements that contain state (sequential)
 - output is function of current and previous inputs
 - state = memory



Components for MIPS subset

- Register,
- Adder
- ALU
- Multiplexer
- Register file
- Program memory
- Data memory
- Bit manipulation components

Components - register

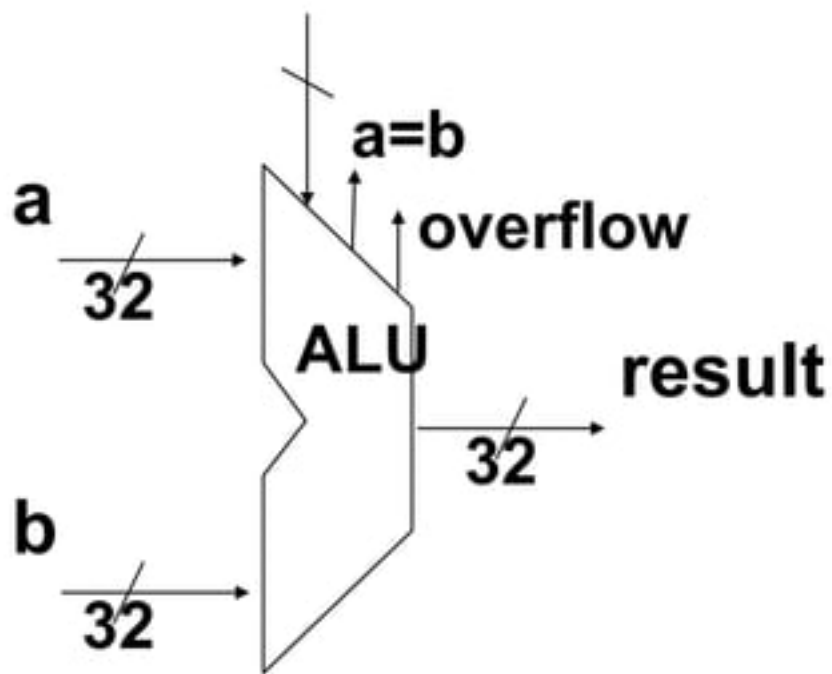


Components - adder



Components - ALU

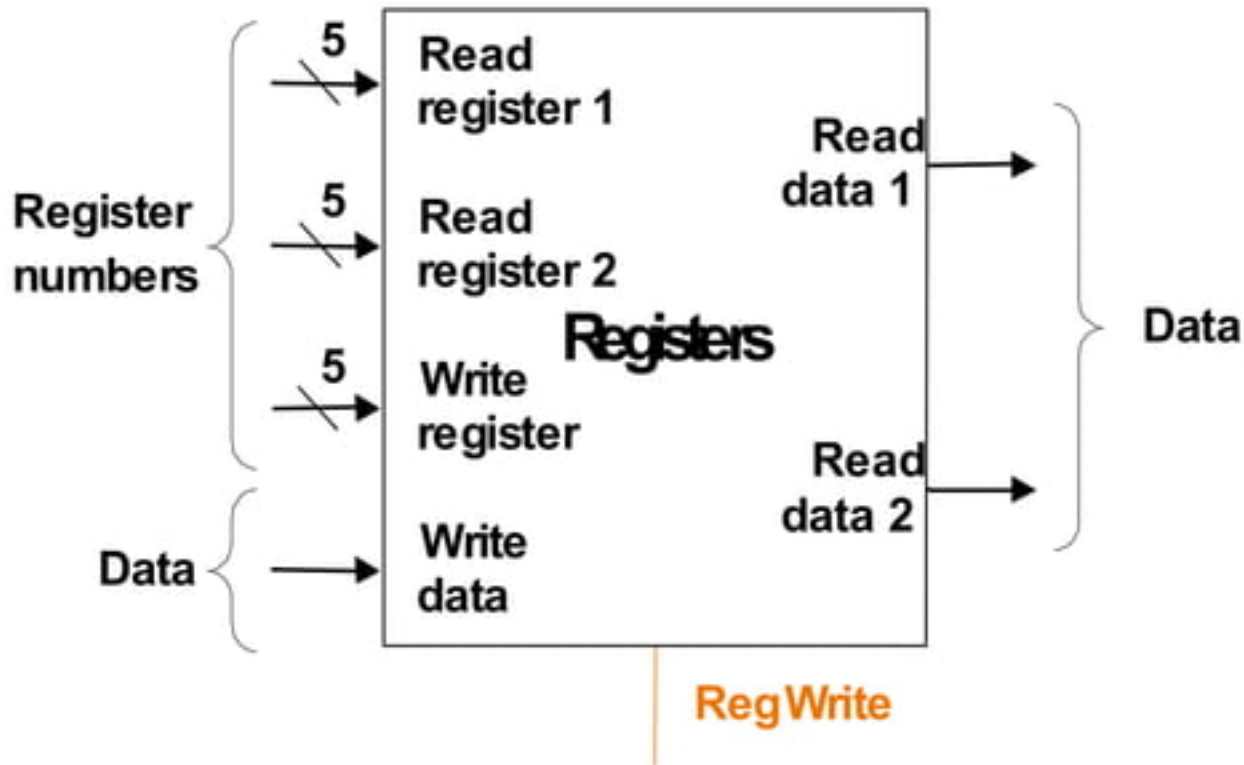
operation




Components - multiplexers




Components - register file







Components - program memory



MIPS components - data memory



Components - bit manipulation circuits



MIPS subset for implementation

- Arithmetic - logic instructions
 - add, sub, and, or, slt
- Memory reference instructions
 - lw, sw
- Control flow instructions
 - beq, j

Datapath for add,sub,and,or,slt

- Fetch instruction
- Address the register file
- Pass operands to ALU actions
- Pass result to register file required
- Increment PC




Format: add \$t0, \$s1, \$s2


000000	10001	10010	01000	000000
--------	-------	-------	-------	--------

100000


Fetching instruction



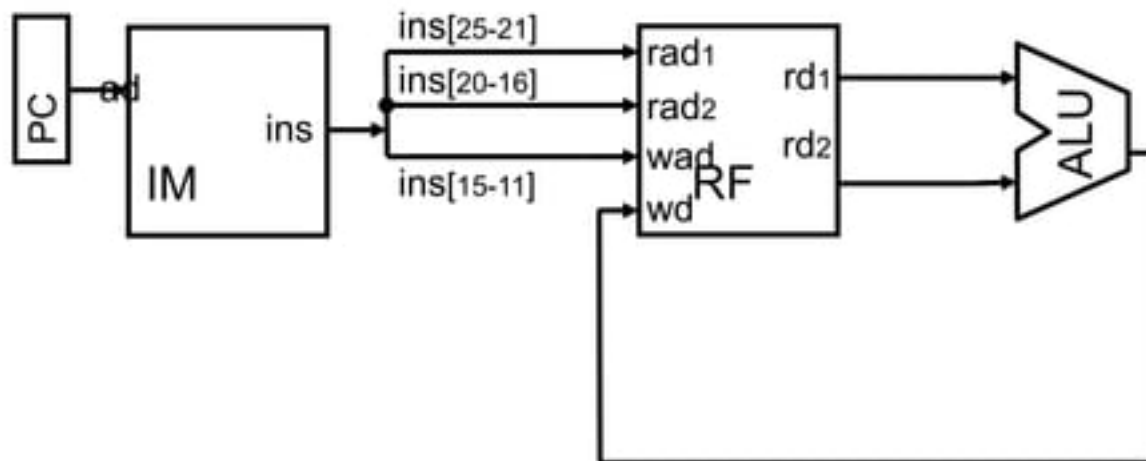
Addressing RF



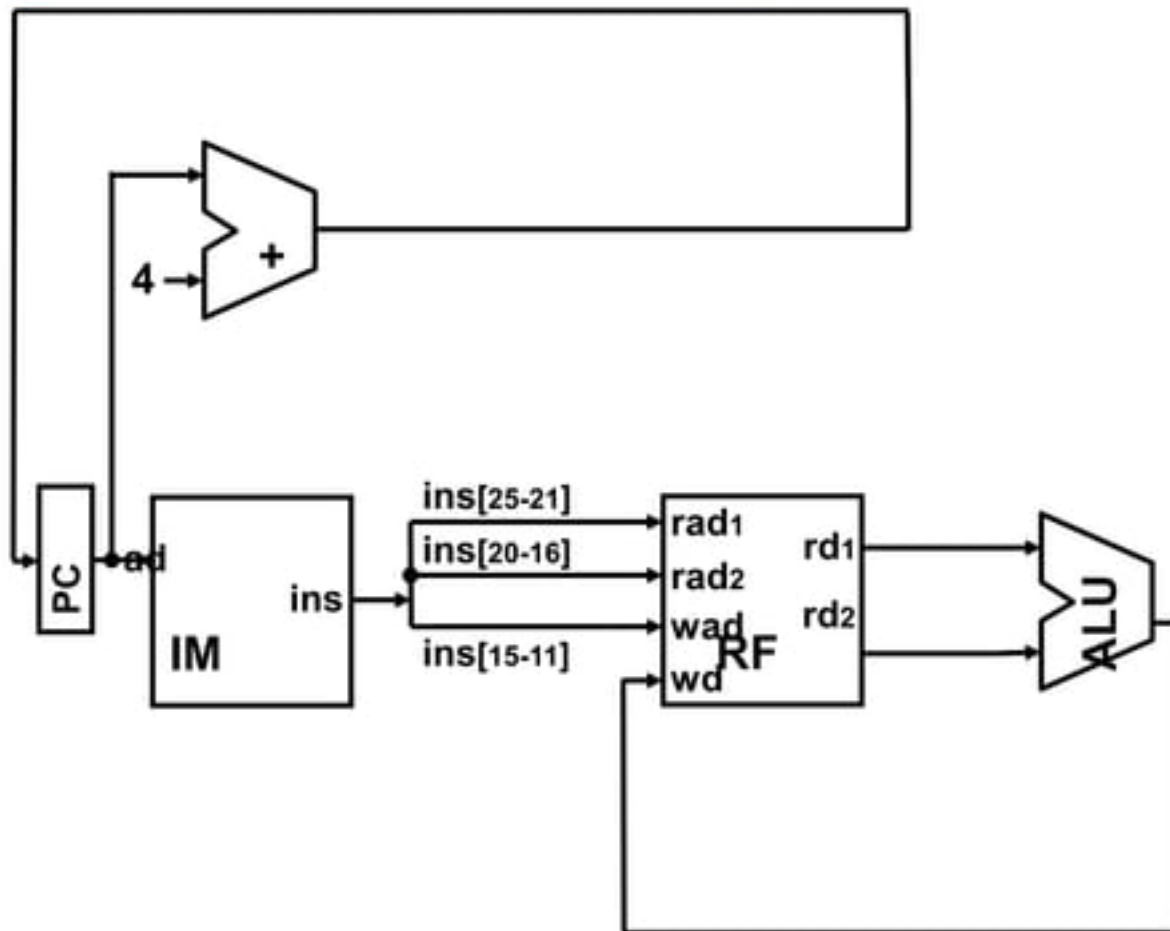
Passing operands to ALU



Passing the result to RF




Incrementing PC

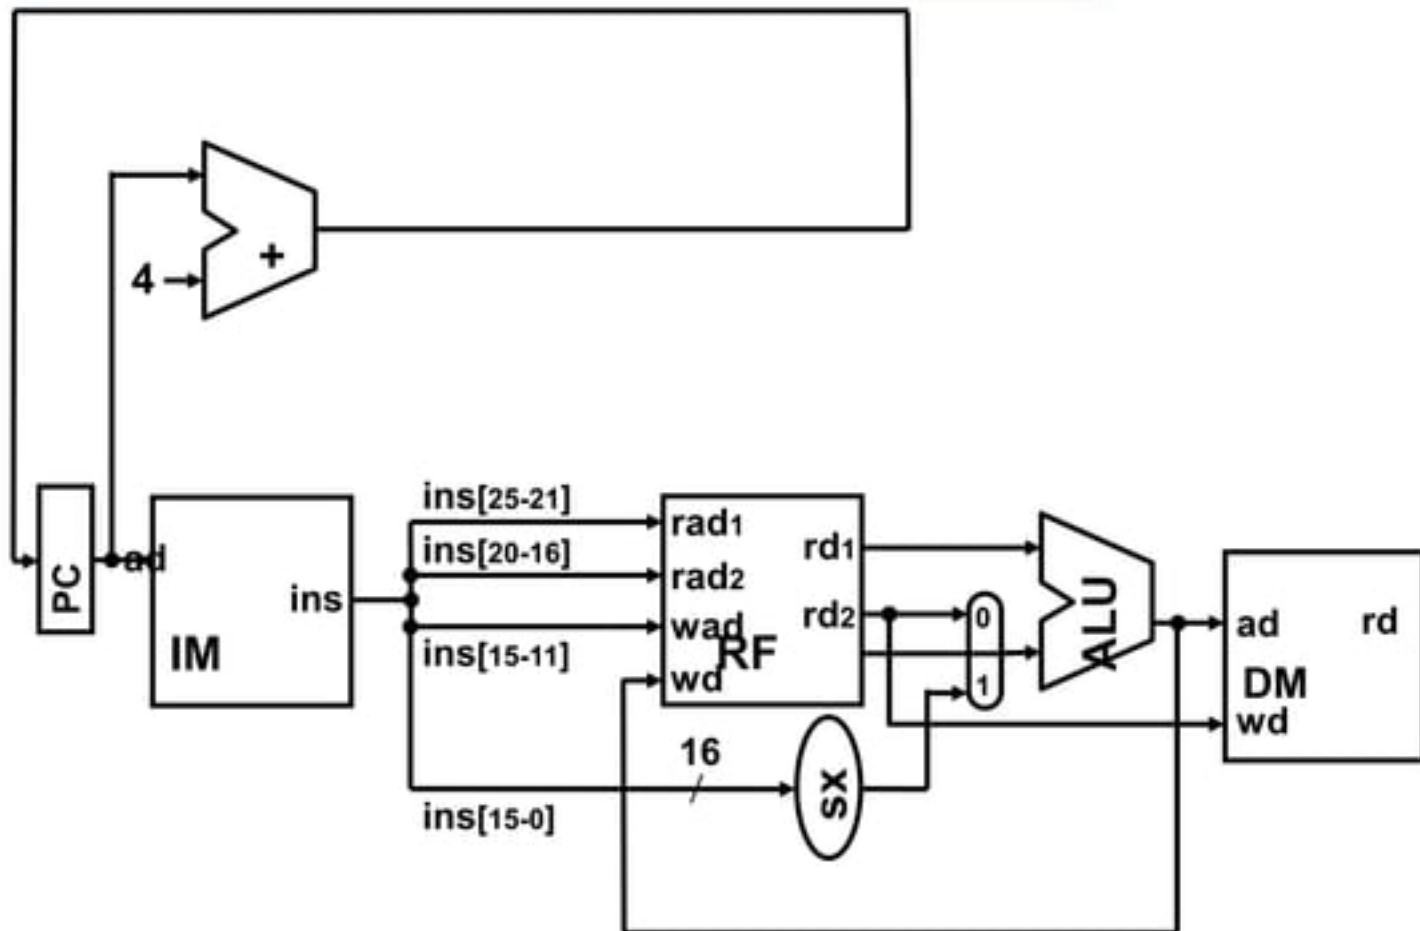


Load and Store instructions

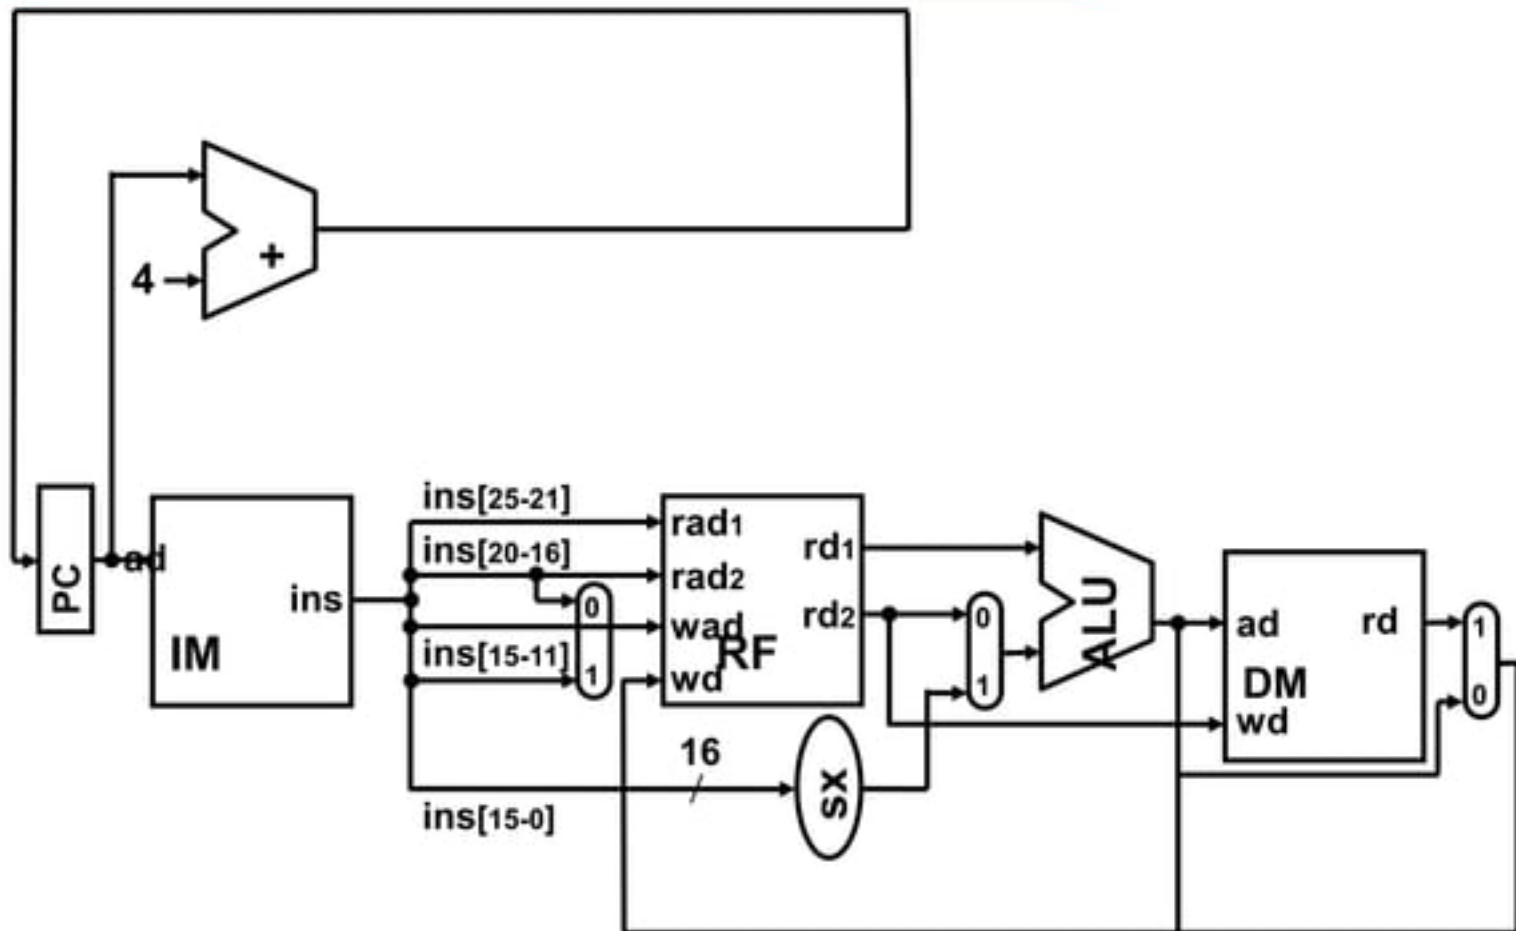
- format : I
- Example: lw \$t0, 32(\$s2)



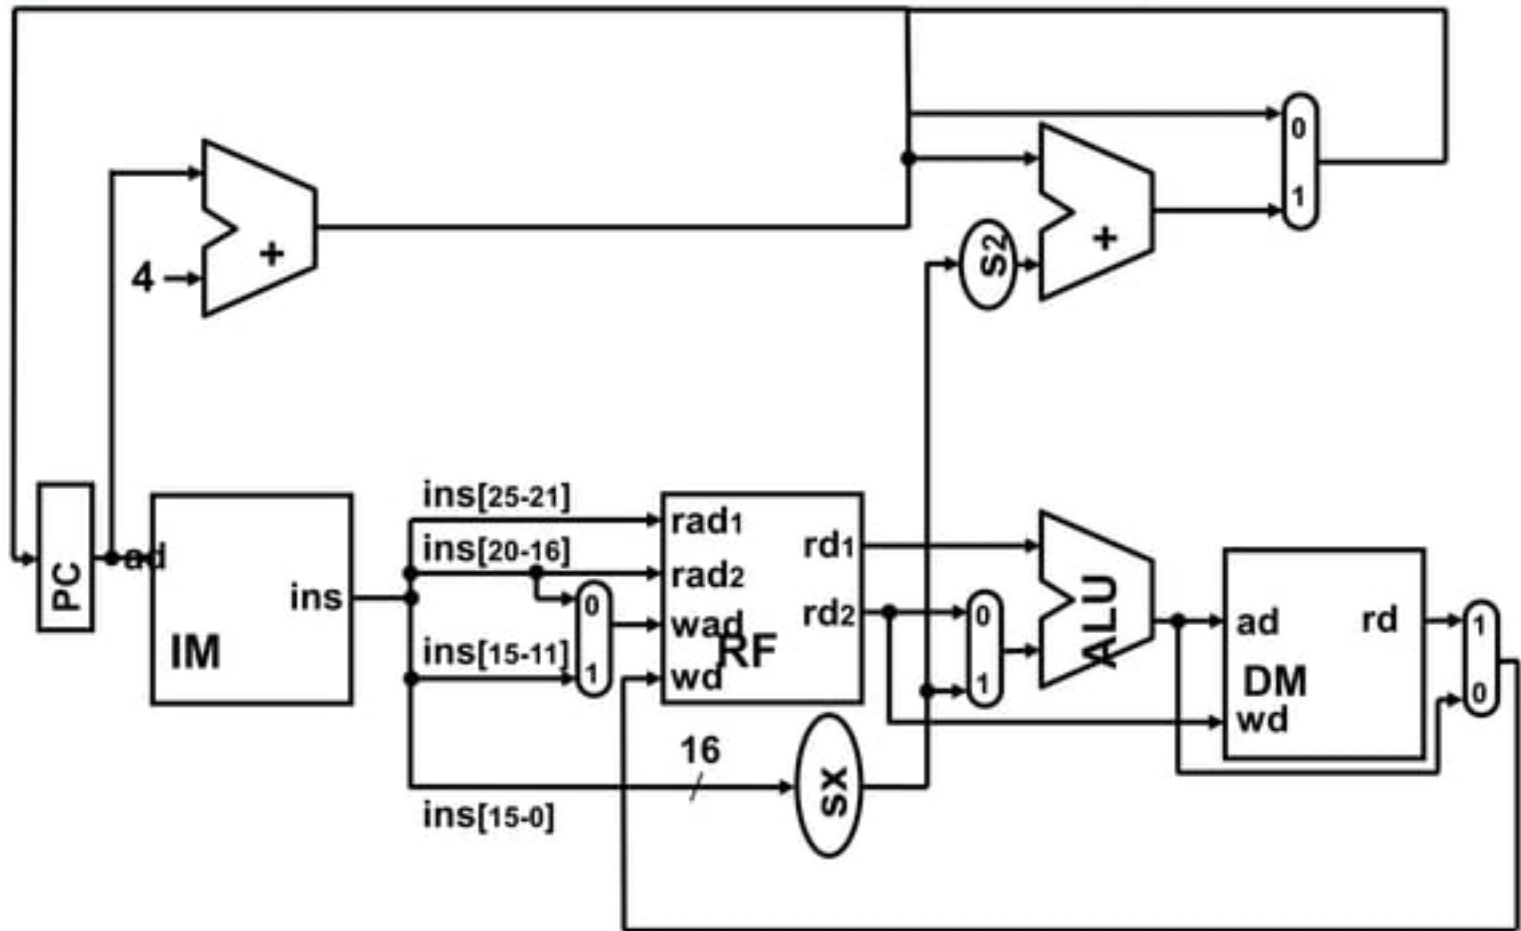
Adding “sw” instruction



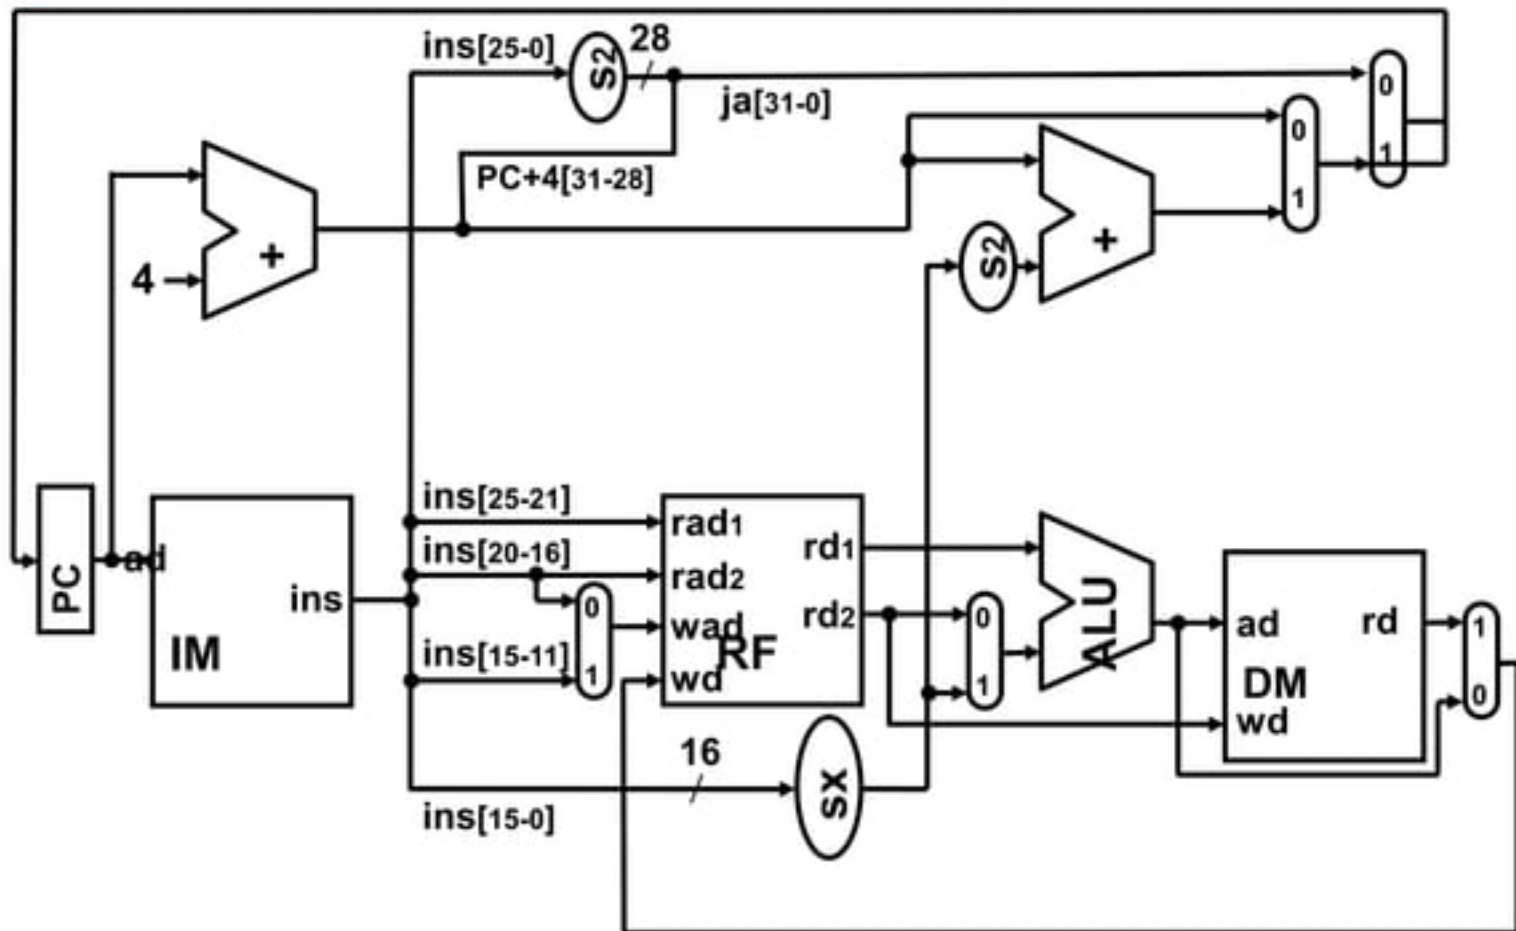
Adding “lw” instruction



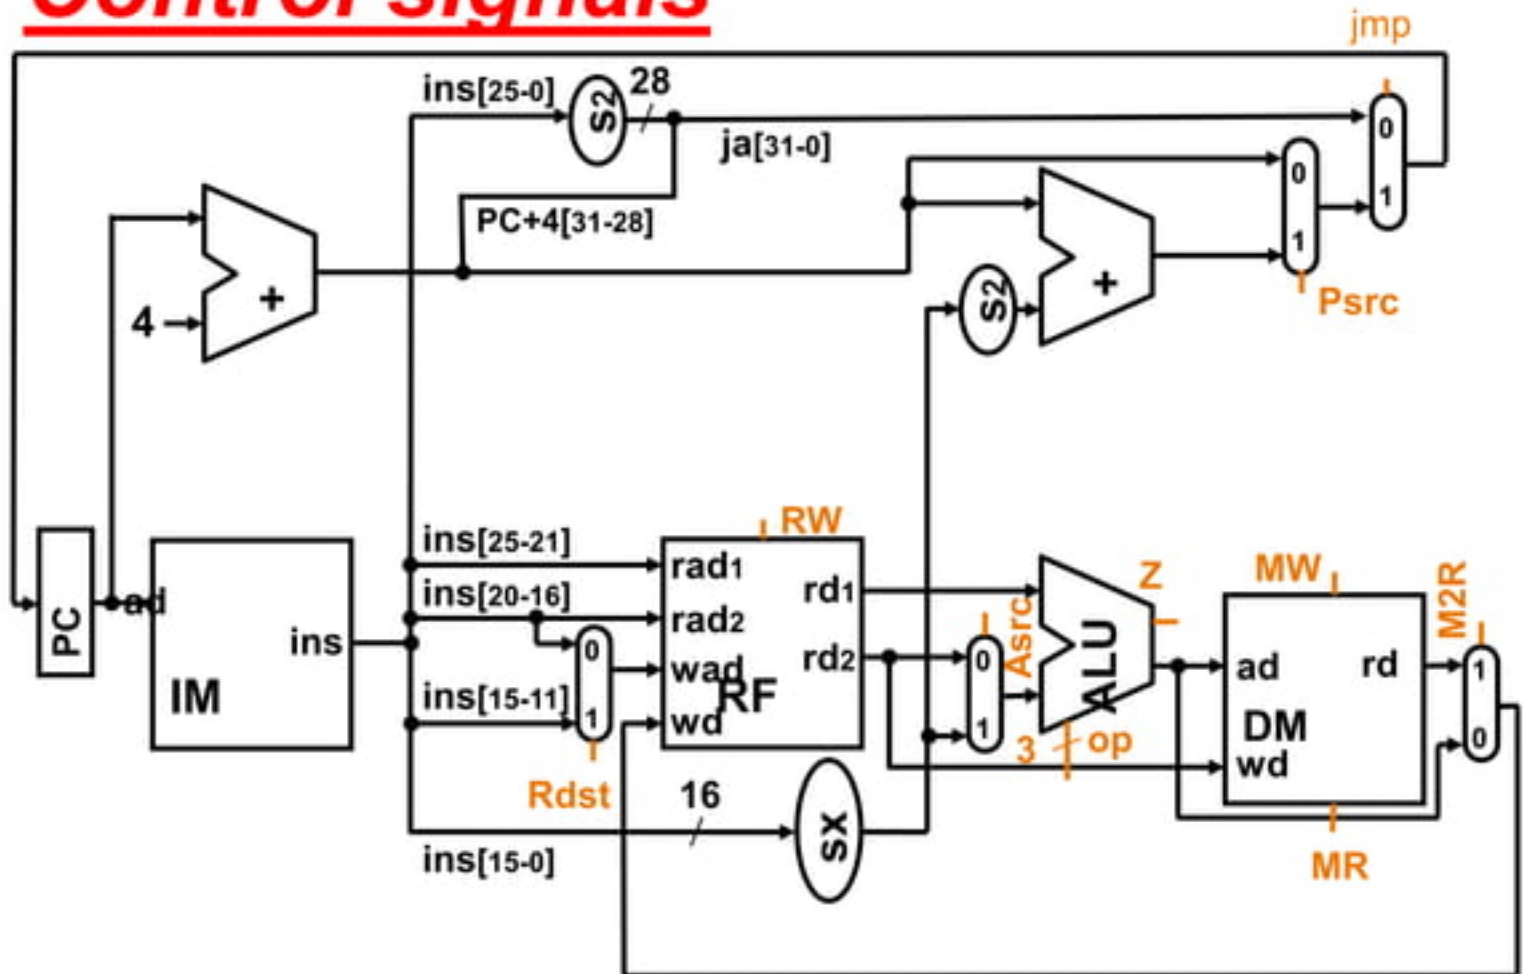
Adding “beq” instruction



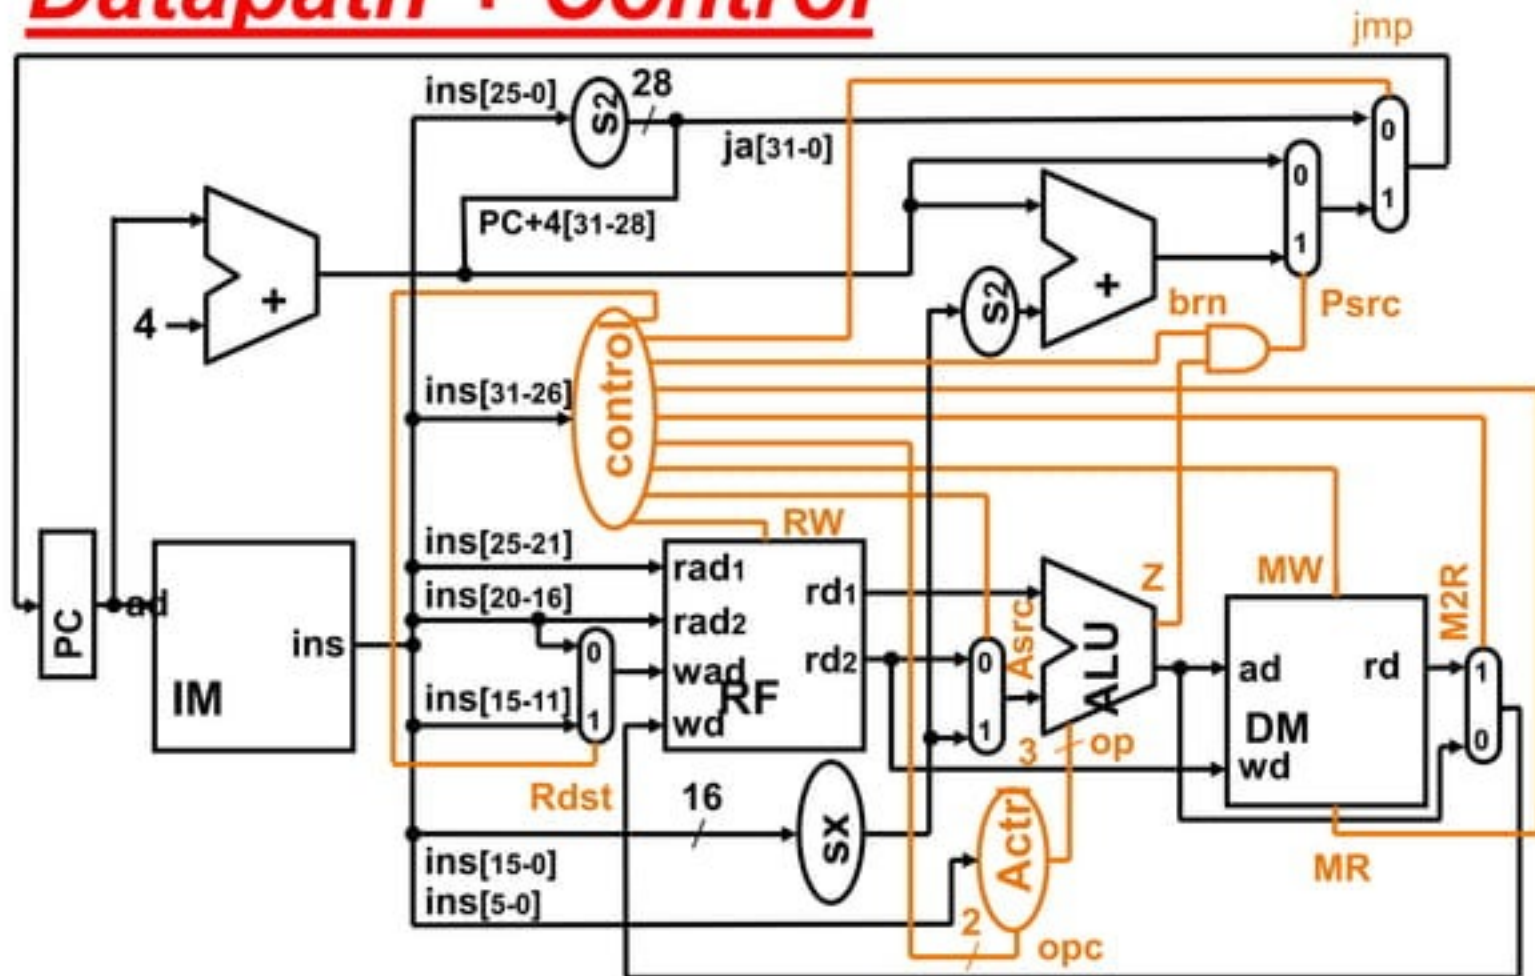
Adding "j" instruction



Control signals



Datapath + Control

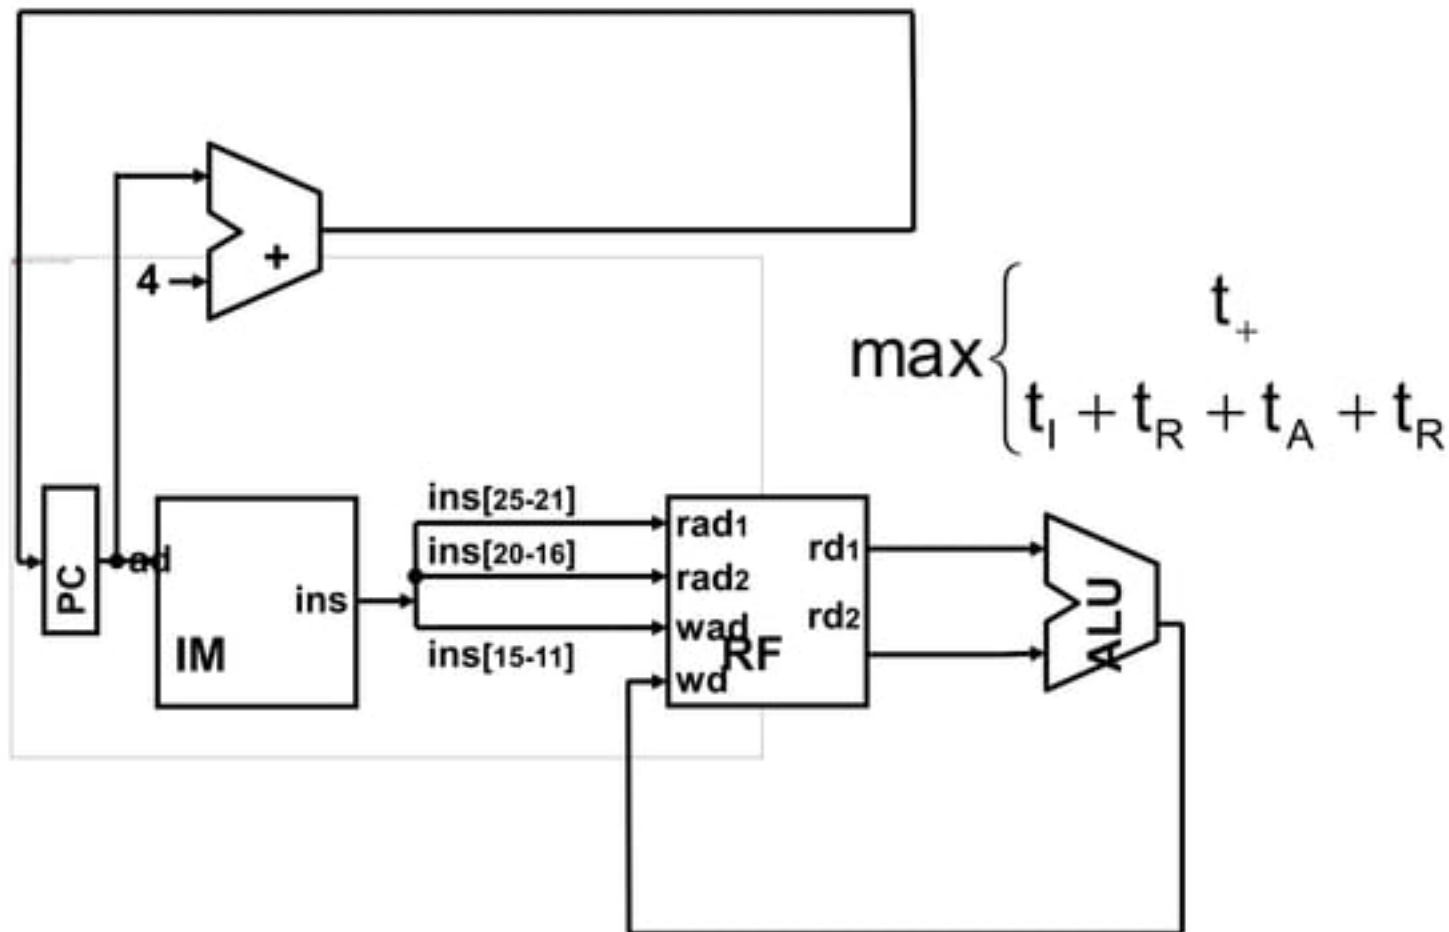


Analyzing performance

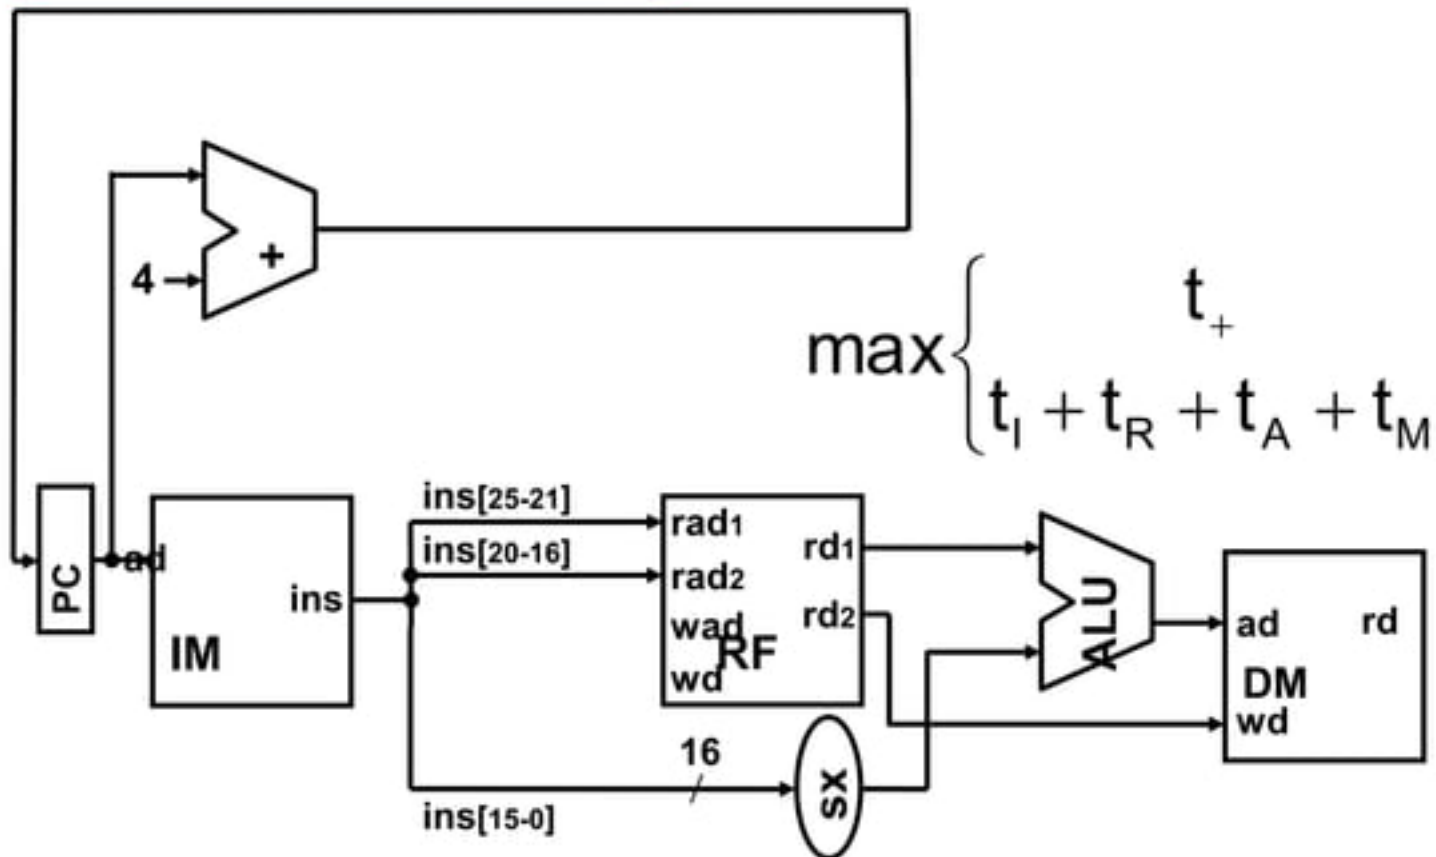
Component delays

• Register		0
• Adder	t_+	
• ALU	t_A	
• Multiplexer	0	
• Register file	t_R	
• Program memory		t_I
• Data memory		t_M
• Bit manipulation components		0

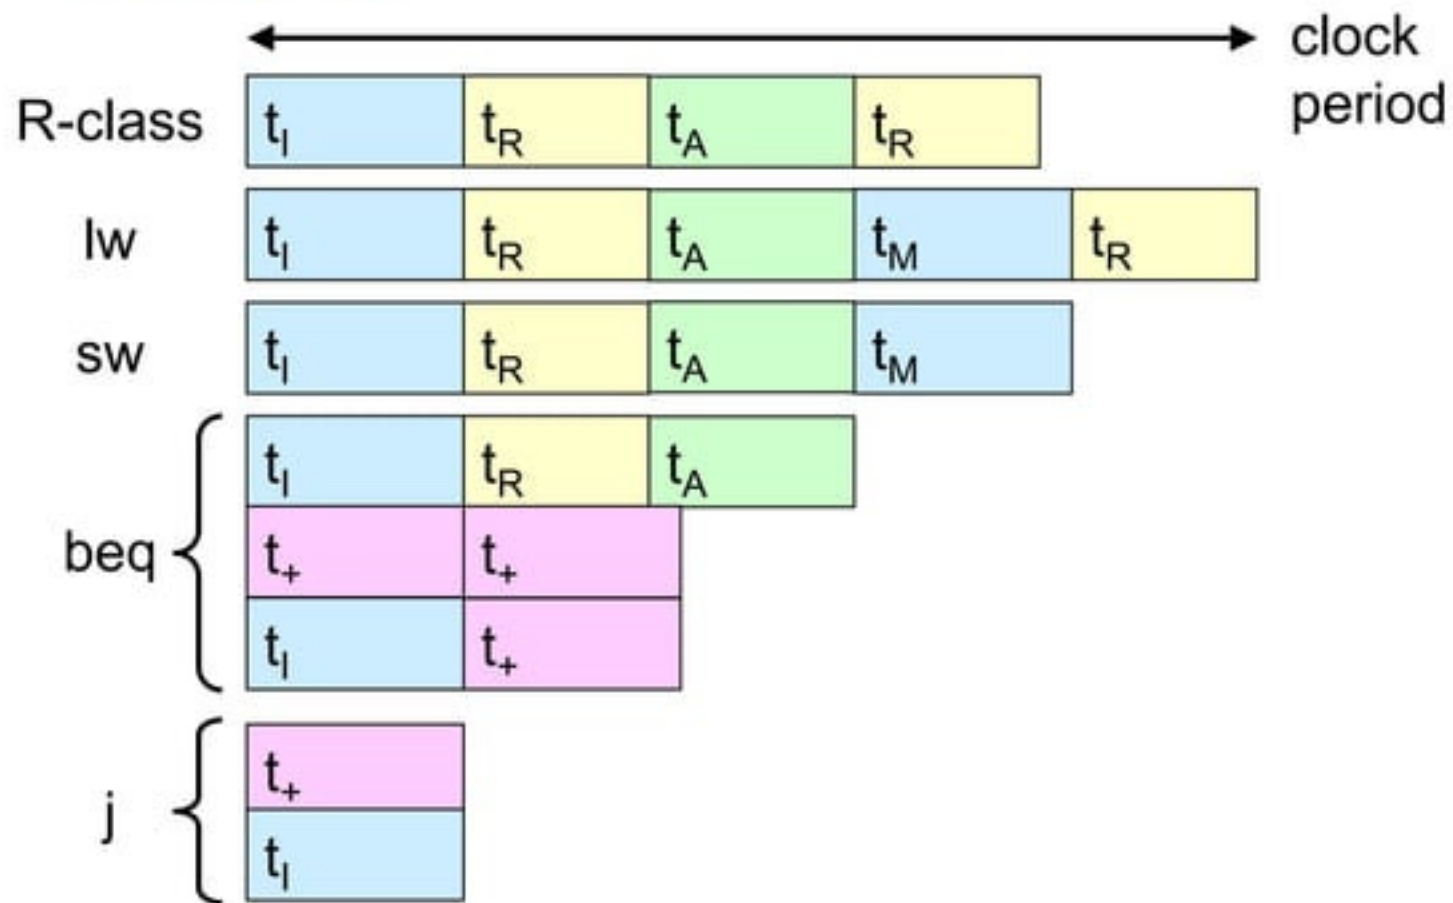
Delay for {add, sub, and, or, slt}



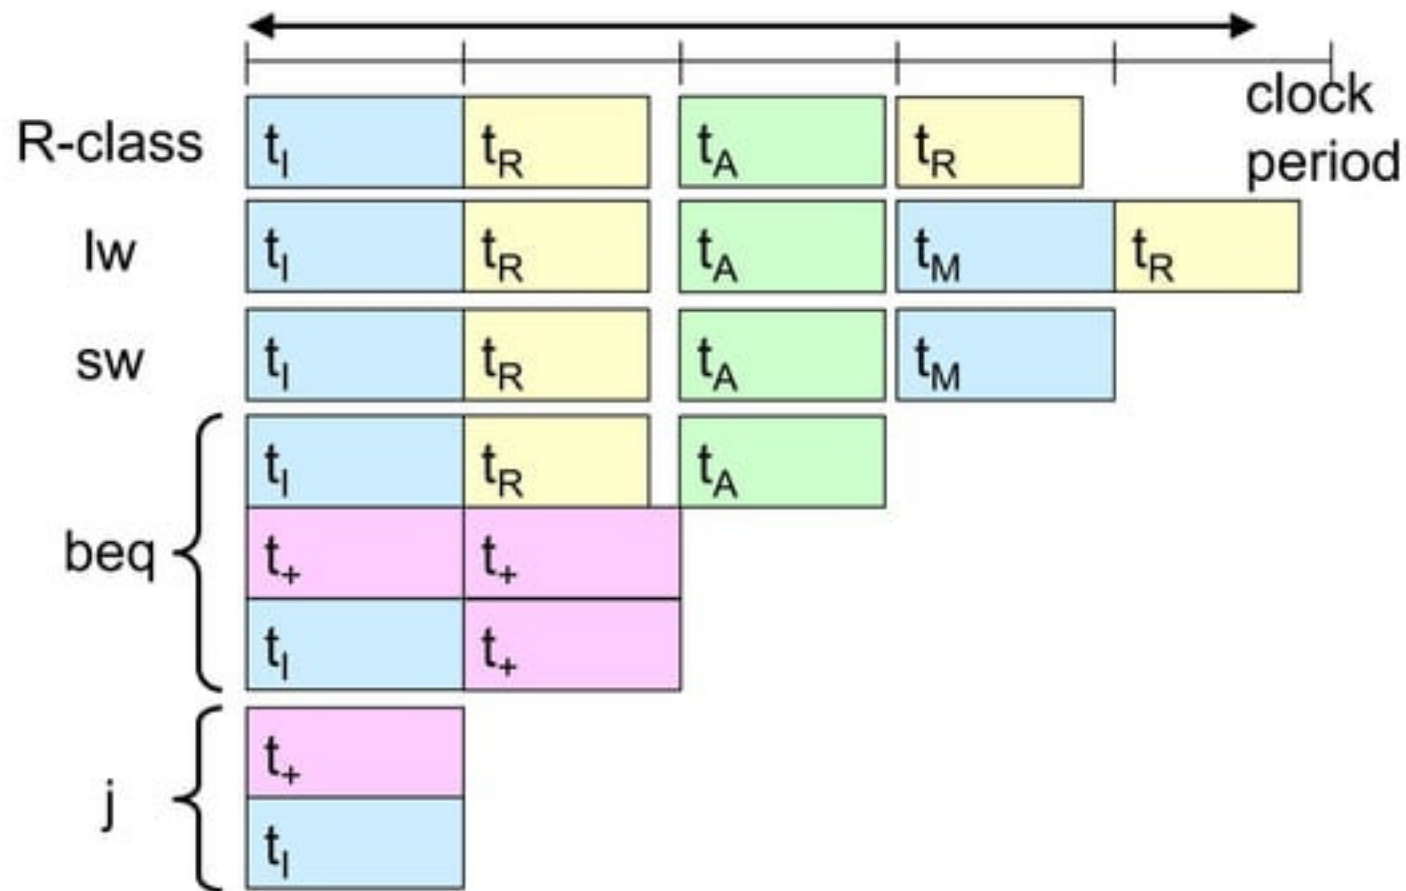
Delay for {sw}




Clock period in single cycle design



Clock period in multi-cycle design



Cycle time and CPI



Pipelined datapath (abstract)

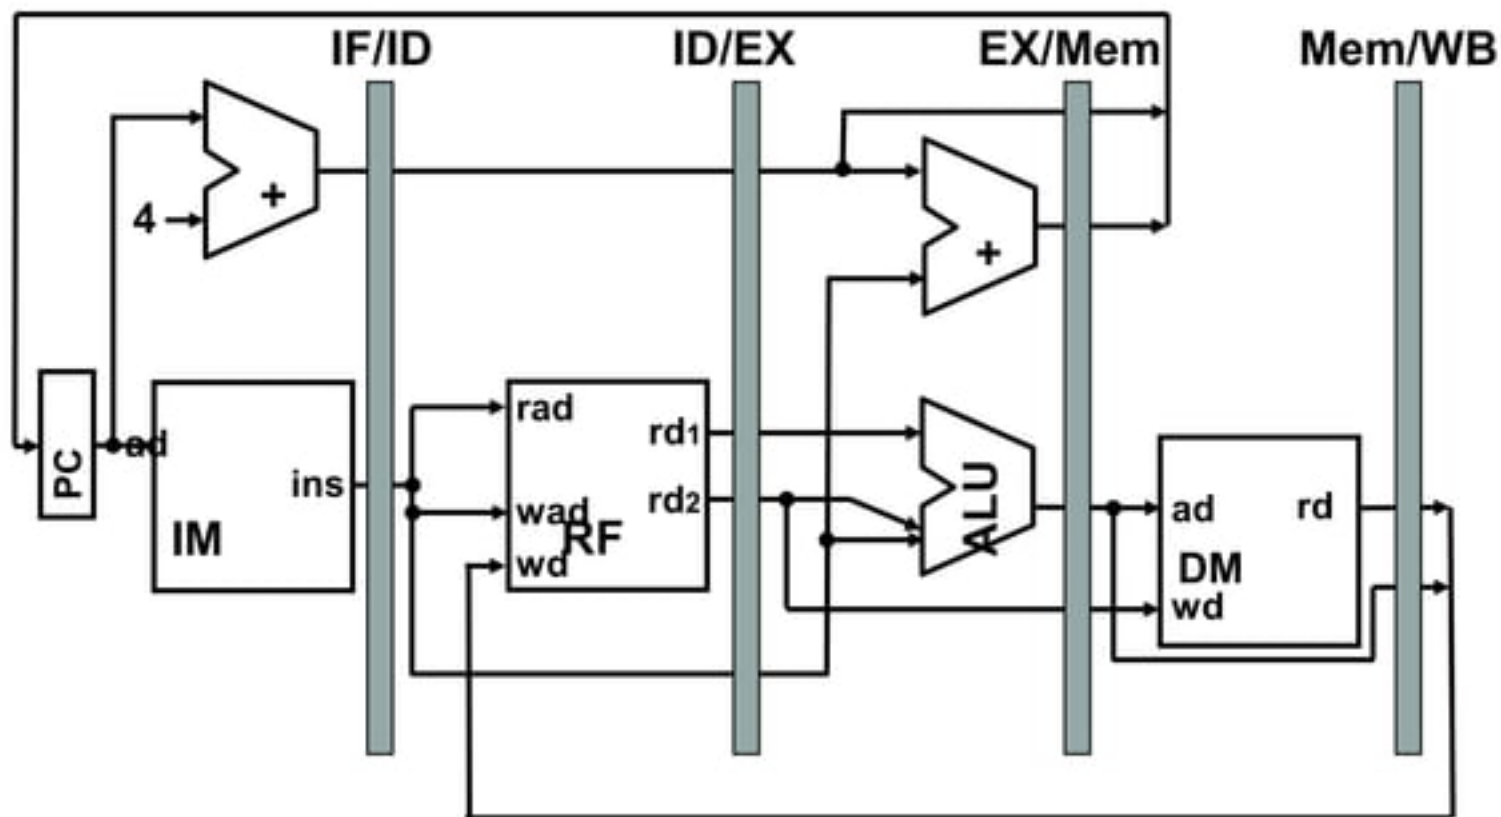
IF

ID


EX

Mem

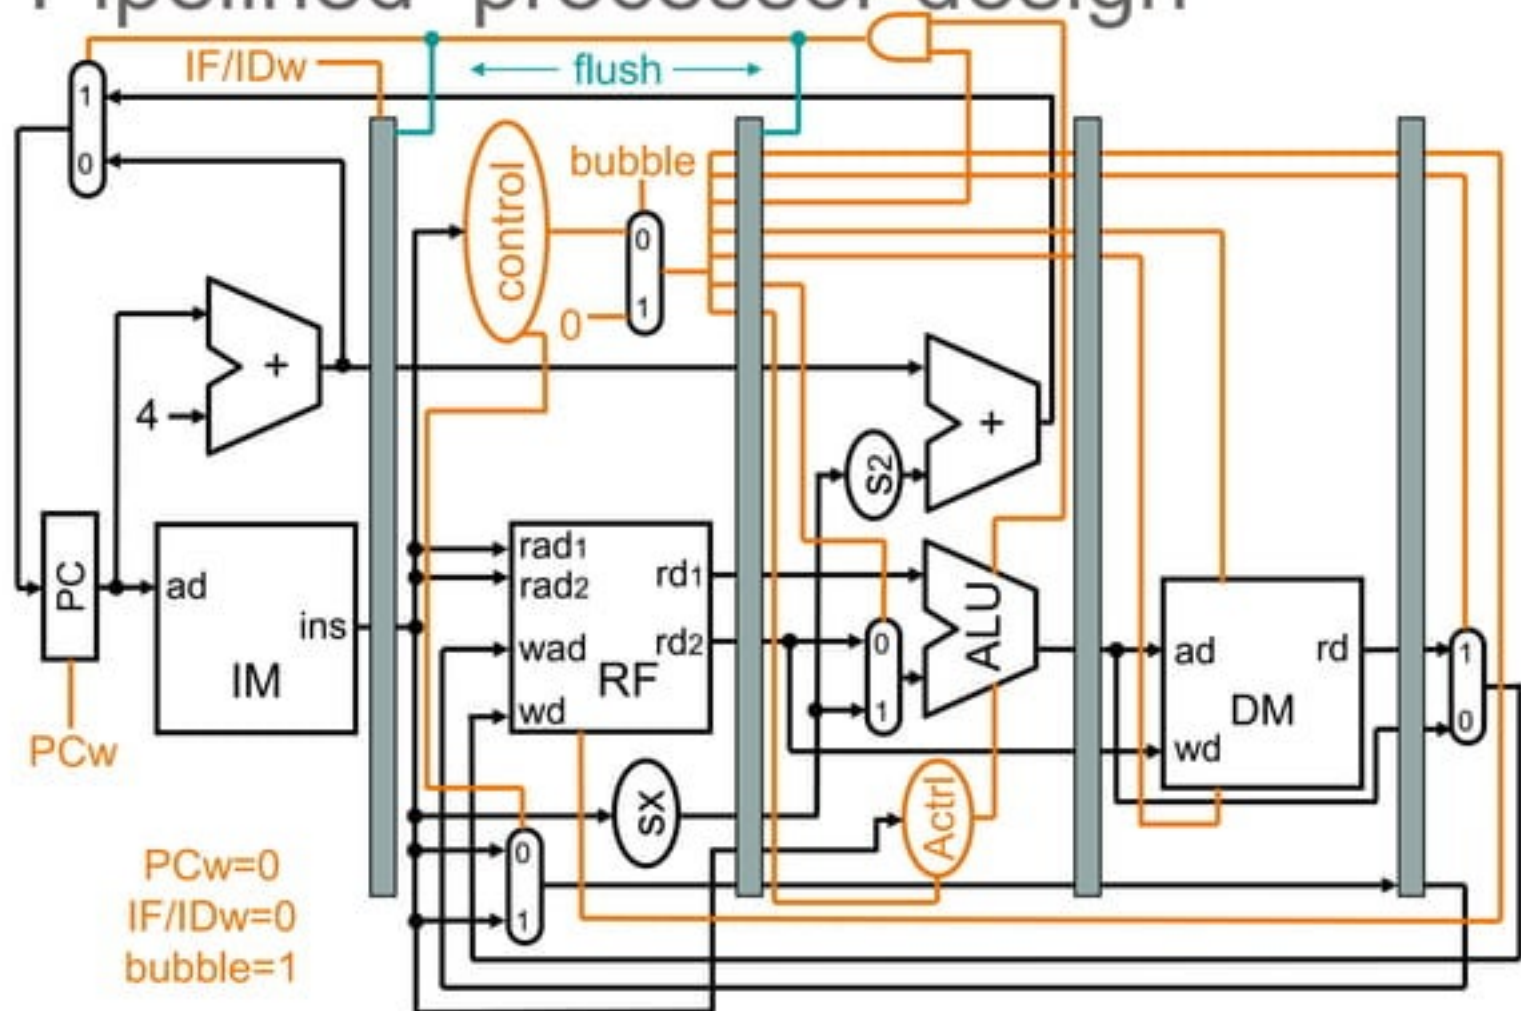
WB



Fetch new instruction every cycle

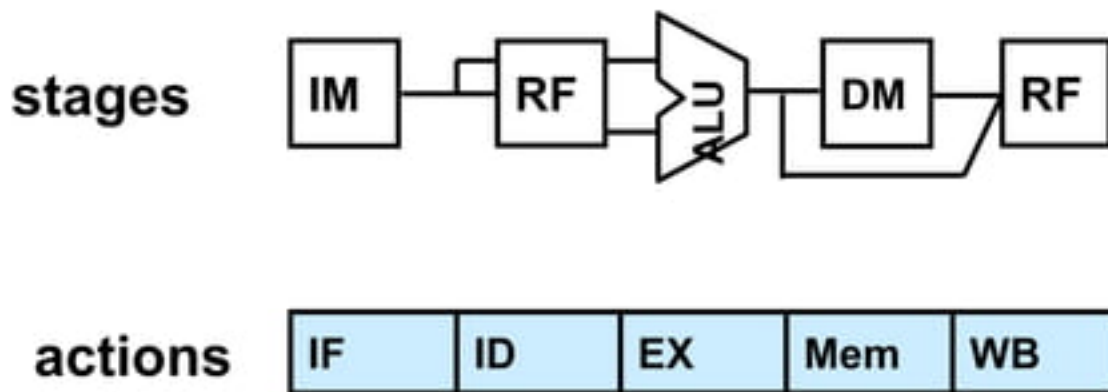


Pipelined processor design




Graphical representation

5 stage pipeline




Usage of stages by instructions



Pipelining

Simple multicycle design :

- Resource sharing across cycles
- All instructions may not take same cycles



- Faster throughput with pipelining


Degree of overlap

Depth


Serial




Overlapped




Pipelined



Shallow



Deep



Hazards in Pipelining

- **Procedural dependencies => Control hazards**
 - **cond and uncond branches, calls/returns**
- **Data dependencies => Data hazards**
 - **RAW (read after write)**
 - **WAR (write after read)**
 - **WAW (write after write)**
- **Resource conflicts => Structural hazards**
 - **use of same resource in different stages**




Data Hazards




Structural Hazards

Caused by Resource Conflicts

- Use of a hardware resource in more than one cycle
- Different sequences of resource usage by different instructions
- Non-pipelined multi-cycle resources




Control Hazards



- the order of *cond eval* and *target addr gen* may be different
- *cond eval* may be done in previous instruction

Pipeline Performance




$$\text{CPI} = 1 + (S - 1) * b$$

$$\text{Time} = \text{CPI} * T / S$$


Improving Branch Performance

- **Branch Elimination**
 - Replace branch with other instructions
- **Branch Speed Up**
 - Reduce time for computing CC and TIF
- **Branch Prediction**
 - Guess the outcome and proceed, undo if necessary
- **Branch Target Capture**

Branch Elimination



Use conditional instructions
(predicated execution)



```
OP1
BC  CC = Z, * + 2
ADD R3, R2, R1
OP2
```




```
OP1
ADD R3, R2, R1, NZ
OP2
```

Branch Speed Up :

Early target address generation

- Assume each instruction is Branch
- Generate target address while decoding
- If target in same page omit translation
- After decoding discard target address if not Branch



Branch Prediction

- Treat conditional branches as unconditional branches / NOP
- Undo if necessary

Strategies:

- Fixed (*always guess inline*)
- Static (*guess on the basis of instruction type*)
- Dynamic (*guess based on recent history*)

Static Branch Prediction

Instr	%	Guess	Branch	Correct
uncond	14.5	always	100%	14.5%
cond	58	never	54%	27%
loop	9.8	always	91%	9%
call/ret	17.7	always	100%	17.7%

Total 68.2%

Branch Target Capture

- Branch Target Buffer (BTB)
- Target Instruction Buffer (TIB)

instr addr

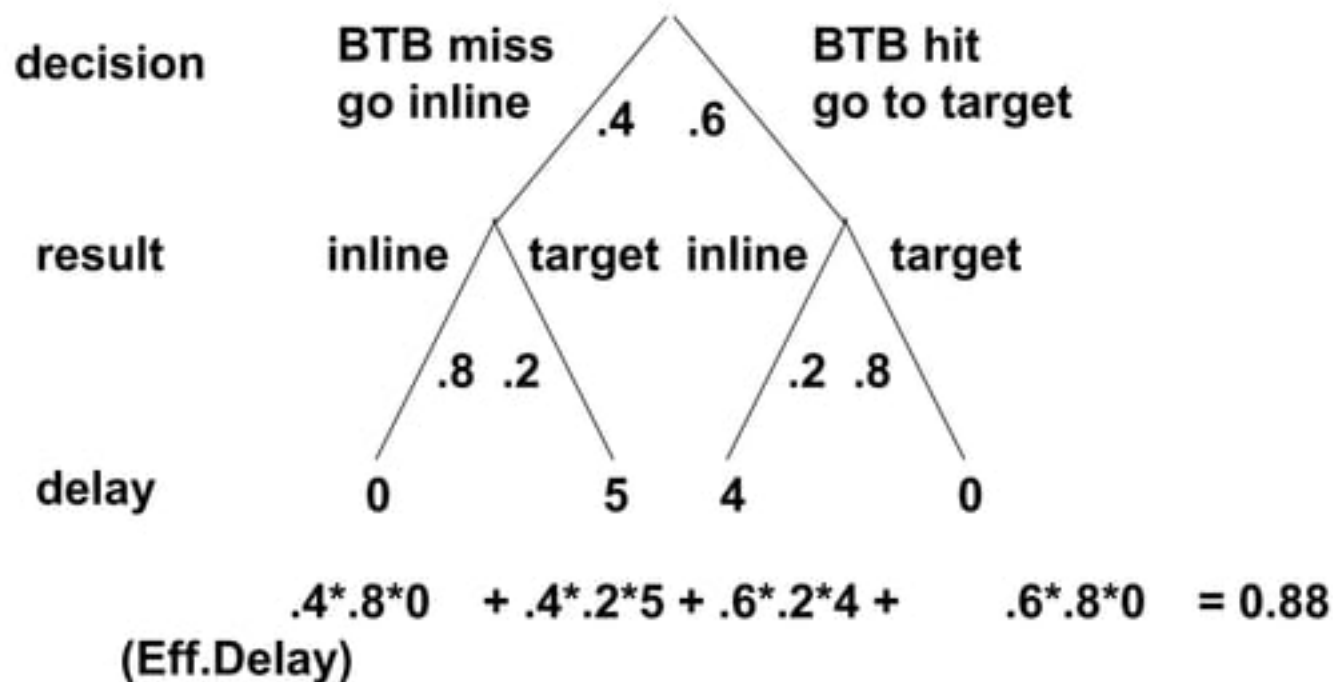
pred stats

target

prob of target change < 5%


target addr
target instr

BTB Performance

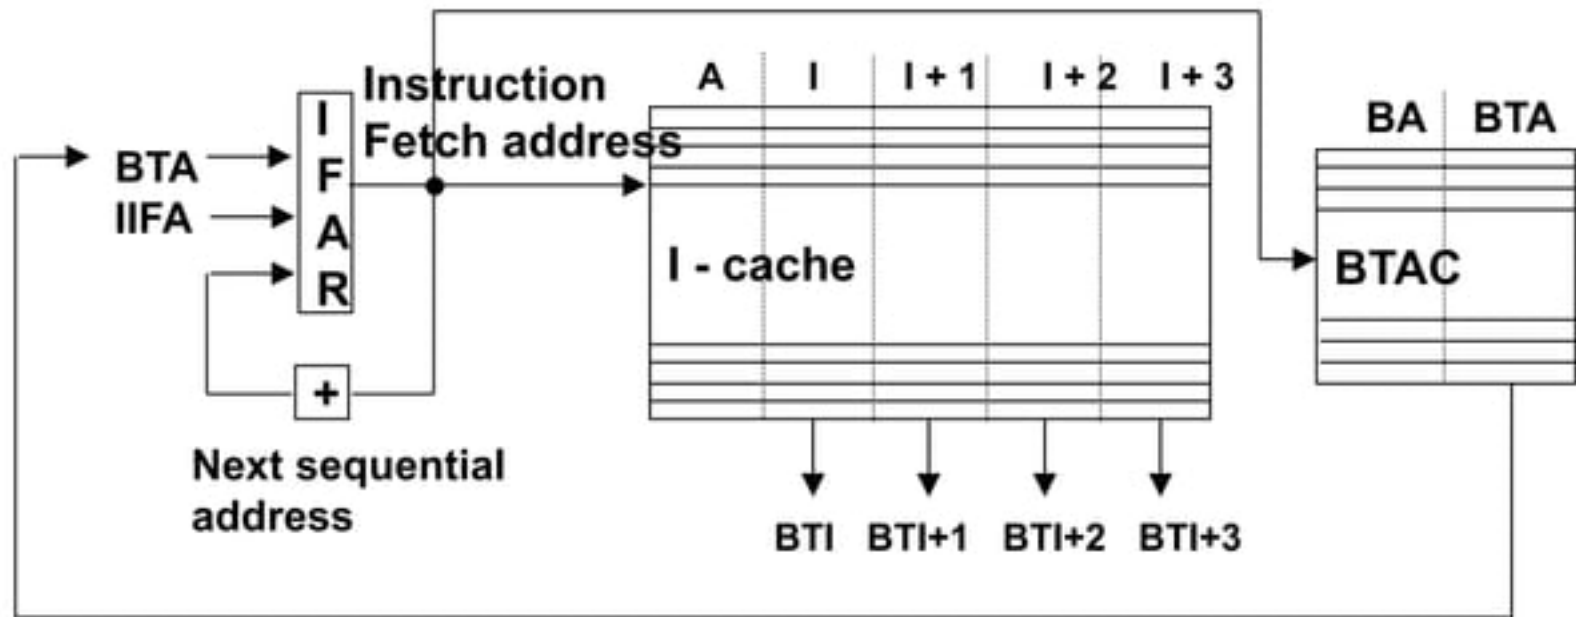


Compute/fetch scheme

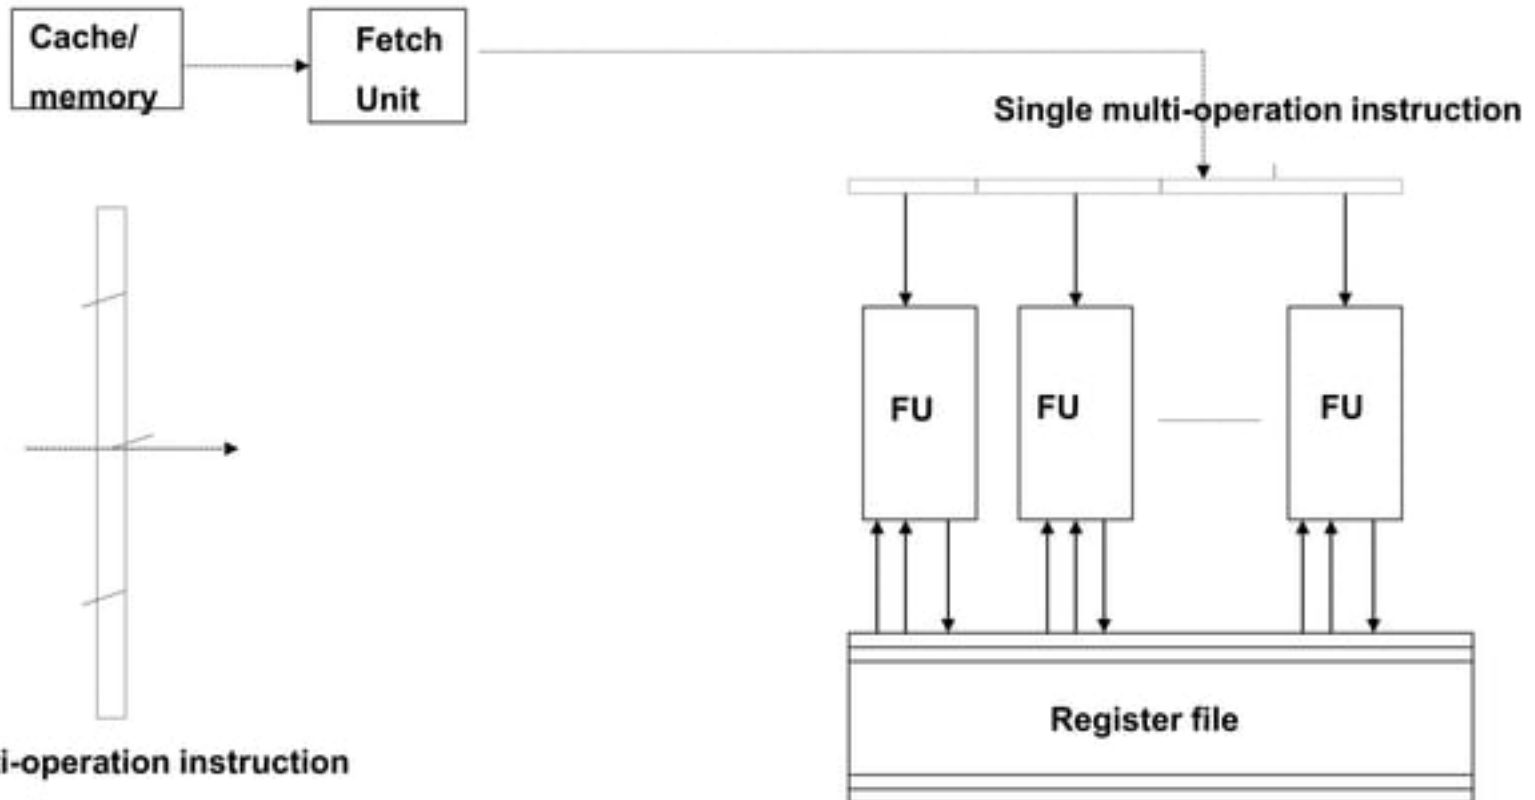
(no dynamic branch prediction)



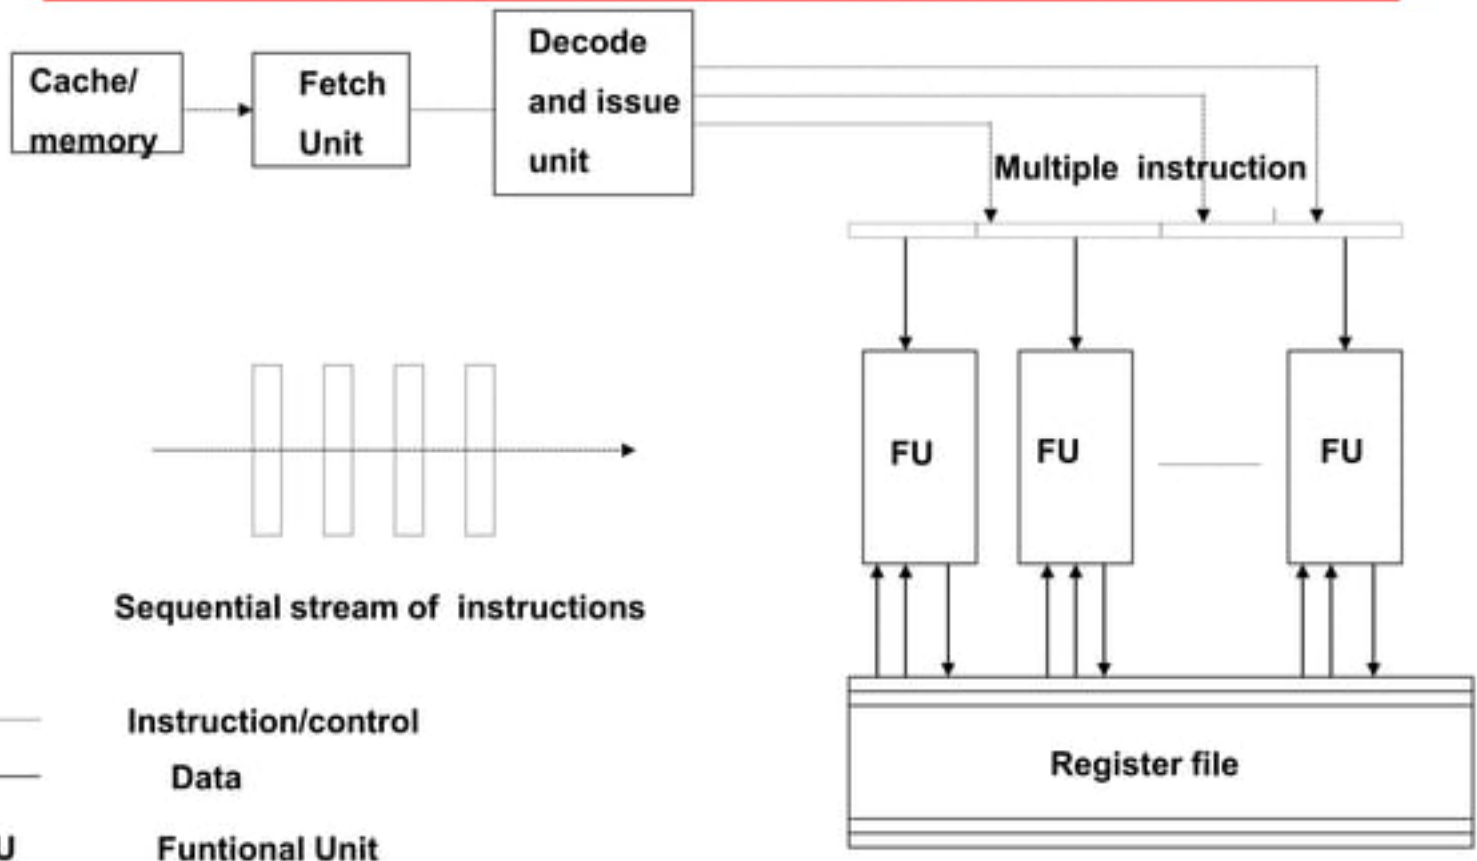
BTAC scheme



ILP in VLIW processors




ILP in Superscalar processors




Why Superscalars are popular *?*

- Binary code compatibility among scalar & superscalar processors of same family
- Same compiler works for all processors (scalars and superscalars) of same family
- Assembly programming of VLIWs is tedious
- Code density in VLIWs is very poor -

Hierarchical structure



Data transfer between levels




Principle of locality & Cache Policies

- Temporal Locality
 - references repeated in time
- Spatial Locality
 - references repeated in space
 - Special case: Sequential Locality

=====

- Read
 - Sequential / Concurrent
 - Simple / Forward
- Load
 - Block load / Load forward / Wrap around
- Replacement
 - LRU / LFU / FIFO / Random

Load policies



Cache miss on AU 1



Block Load



Load Forward



**Fetch Bypass
(wrap around
load)**



Fetch Policies

- Demand fetching
 - fetch only when required (miss)
 - Hardware prefetching
 - automatically prefetch next block
 - Software prefetching
 - programmer decides to prefetch
- questions:
- how much ahead (prefetch distance)
 - how often

Write Policies

- Write Hit
 - Write Back
 - Write Through
- Write Miss
 - Write Back
 - Write Through
 - With Write Allocate
 - With No Write Allocate

Cache Types

Instruction | Data | Unified | Split

Split vs. Unified:

- Split allows specializing each part
- Unified allows best use of the capacity

On-chip | Off-chip

- on-chip : fast but small
- off-chip : large but slow

References

1. Patterson, D A.; Hennessy, J L. *Computer Organization and Design: The Hardware/software Interface*. Morgan Kaufman, 2000
2. Sima, T, FOUNTAIN, P KACSUK, *Advanced Computer Architectures: A Design Space Approach*, Pearson Education, 1998
3. Flynn M J, *Computer Architecture: Pipelined and Parallel Processor Design*, Narosa publishing India, 1999
4. John L. Hennessy, David A. Patterson, *Computer architecture: a quantitative approach*, 2nd Ed, Morgan Kauffman, 2001

Thanks